

# MNIST DIGIT CLASSIFIER - REPORT

## Key Findings:

The CNN-based classifier demonstrated excellent performance on the MNIST dataset, achieving over 99.4% training accuracy within just 5 epochs. The training loss decreased steadily, indicating effective learning. The confusion matrix indicates that most digits were classified correctly, with very few misclassifications.

## Model Architecture and Design Choices:

The implemented model is a Convolutional Neural Network (CNN), which is well-suited for image data. The architecture includes:

- Two convolutional layers for local feature extraction.
- ReLU activations to introduce non-linearity.
- MaxPooling layers to reduce dimensionality.
- Fully connected layers to classify extracted features into 10 digit classes.

This design choice significantly boosts performance compared to a simple Feedforward Neural Network due to CNN's ability to learn spatial hierarchies in images.

## Confusion Matrix Analysis:

The confusion matrix shows high classification accuracy across all digits. Some observed misclassifications:

- Digit '5' is occasionally confused with '3' and '8', possibly due to similar curvature in handwritten forms.
- Digit '9' sometimes gets confused with '4', likely due to sharp strokes in certain handwriting styles.
- Digits '7' and '1' have minor confusion, which is common due to similar vertical structures. These misclassifications are minimal and expected given handwriting variability.

## Training Plots:

